

FINAL PROJECT

MUSIC DATA ANALYSIS

Table of Contents

Section – 1 - Project Overview	2
1.1 Fields present in the data files	2
1.2 LookUp Tables	2
1.3 DATASET	3
1.4 Data Enrichment	4
1.5 Data Analysis	4
1.6 Challenges and Optimizations	4
1.7 Flow of operations	5
Section -2 – Design of the Project	5
2.1 Low Level Design	5
2.2 High Level Design	6
Section-3-Hadoop Eco-System Implementation	6
Section-4 –Data Ingestion, Formatting, Enrichment and Filtering	8
4.1 Stage – 1 – Data Ingestion	8
4.2 Stage – 2 - Data Formatting	15
4.3 Stage – 3 - Data Enrichment & Filtering	18
4.4 Stage – 4 - Data Analysis using Spark	22
4.5 Stage – 5 – Data Storage in MYSQL	27
Job Scheduling	30
Highlights of the Project	31
Project End Conclusion	32

SECTION – 1

A leading music-catering company is planning to analyse large amount of data received from varieties of sources, namely mobile app and website to track the behaviour of users, classify users, calculate royalties associated with the song and make appropriate business strategies. The file server receives data files periodically after every 3 hours.

1.1 Fields present in the data files

Data files contain below fields.

Column Name/Field Name	Column Description/Field Description
User_id	Unique identifier of every user
Song_id	Unique identifier of every song
Artist_id	Unique identifier of the lead artist of the song
Timestamp	Timestamp when the record was generated
Start_ts	Start timestamp when the song started to play
End_ts	End timestamp when the song was stopped
Geo_cd	Can be 'A' for USA region, 'AP' for asia pacific region, 'J' for Japan region, 'E' for europe and 'AU' for australia region
Station_id	Unique identifier of the station from where the song was played
Song_end_type	How the song was terminated. 0 means completed successfully 1 means song was skipped 2 means song was paused 3 means other type of failure like device issue, network error etc.
Like	0 means song was not liked 1 means song was liked
Dislike	0 means song was not disliked 1 means song was disliked

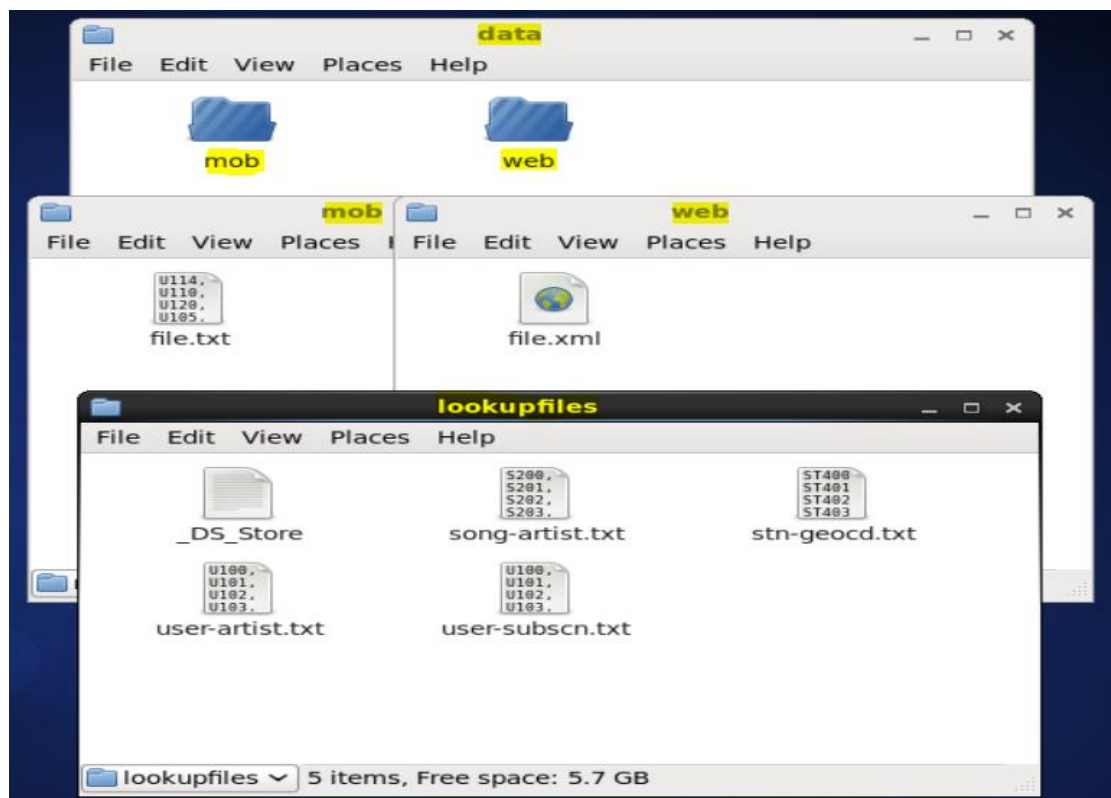
1.2 LookUp Tables

There is some existing look up tables present in NoSQL databases. They play an important role in data enrichment and analysis.

Table Name	Description
Station_Geo_Map	Contains mapping of a geo_cd with station_id
Subscribed_Users	Contains user_id, subscription_start_date and subscription_end_date. Contains details only for subscribed users
Song_Artist_Map	Contains mapping of song_id with artist_id alongwith royalty associated with each play of the song
User_Artist_Map	Contains an array of artist_id(s) followed by a user_id

1.3 Dataset

1. Data coming from web applications reside in /data/web and has xml format.
2. Data coming from mobile applications reside in /data/mob and has csv format.
3. Data present in lookup directory should be used in HBase.



1.4 Data Enrichment

Rules for data enrichment,

1. If any of like or dislike is NULL or absent, consider it as 0.
2. If fields like **Geo_cd** and **Artist_id** are NULL or absent, consult the lookup tables for fields **Station_id** and **Song_id** respectively to get the values of **Geo_cd** and **Artist_id**.
3. If corresponding lookup entry is not found, consider that record to be invalid

NULL or absent field	Look up field	Look up table (Table from which record can be updated)
Geo_cd	Station_id	Station_Geo_Map
Artist_id	Song_id	Song_Artist_Map

1.5 Data Analysis

It is not only the data which is important, rather it is the insight it can be used to generate important. Once we have made the data ready for analysis, we have to perform below analysis on a daily basis.

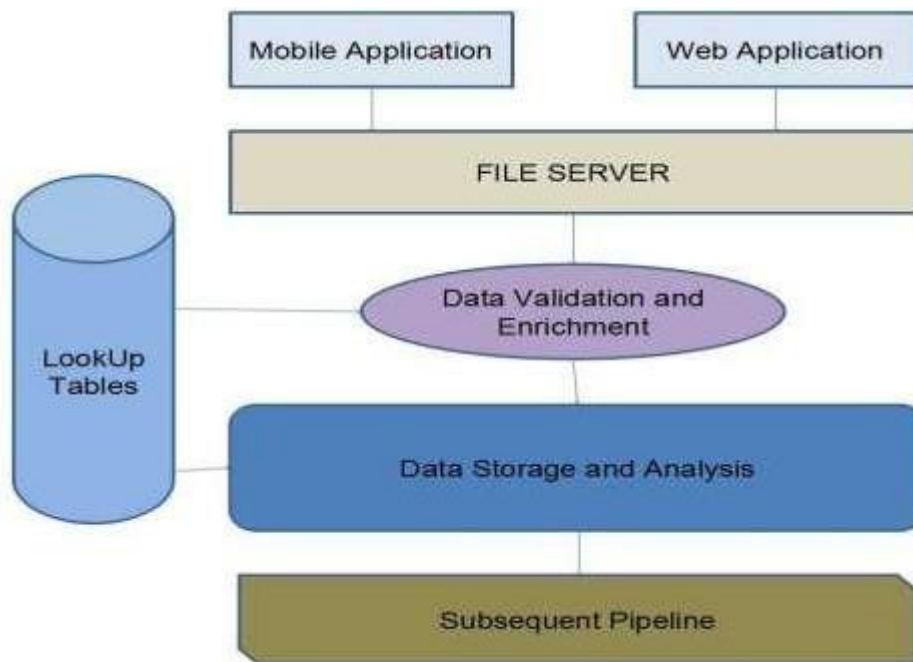
1. Determine top 10 **station_id**(s) where maximum number of songs were played, which were liked by unique users.
2. Determine total duration of songs played by each type of user, where type of user can be '**subscribed**' or '**unsubscribed**'. An unsubscribed user is the one whose record is either not present in **Subscribed_users** lookup table or has **subscription_end_date** earlier than the timestamp of the song played by him.
3. Determine top 10 connected artists. Connected artists are those whose songs are most listened by the unique users who follow them.
4. Determine top 10 songs who have generated the maximum revenue. Royalty applies to a song only if it was liked or was completed successfully or both.
5. Determine top 10 unsubscribed users who listened to the songs for the longest duration.

1.6 Challenges and Optimizations

1. LookUp tables are in NoSQL databases. Integrate them with the actual data flow.
2. Try to make joins as less expensive as possible.
3. Data Cleaning, Validation, Enrichment, Analysis and Post Analysis have to be automated. Try using schedulers.
4. Appropriate logs have to maintain to track the behaviour and overcome failures in the pipeline.

BIG DATA PROJECT

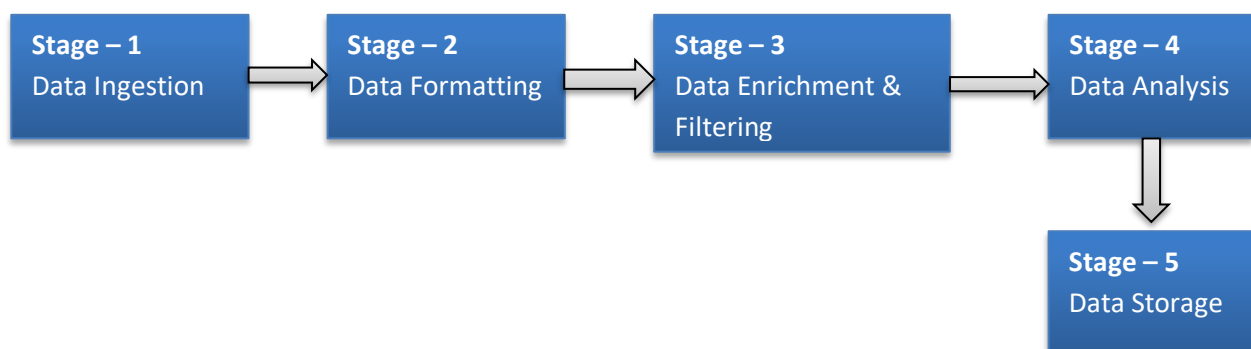
1.7 Flow of operation



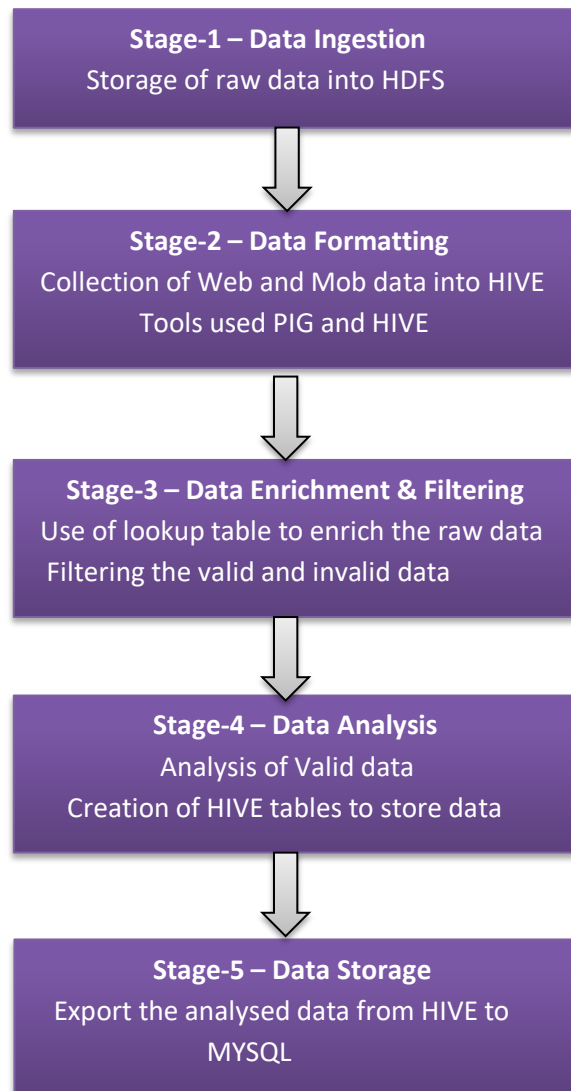
This figure gives an overview of the project

SECTION - 2 – Design of Project

2.1 Low level design



2.2 High level Design



SECTION – 3 – Hadoop Eco-System Implementation

1. We have created a batch file “**start-daemon.sh**” which starts the daemons such as **hive**, **hbase**, **Mysql** and rest of the all **hadoop** daemons.

```
1  #!/bin/bash
2
3  if [ -f "/home/acadgild/project/logs/current-batch.txt" ]
4  then
5      echo "Batch File Found!"
6  else
7      echo -n "1" > "/home/acadgild/project/logs/current-batch.txt"
8  fi
9
10 chmod 775 /home/acadgild/project/logs/current-batch.txt
11 batchid=`cat /home/acadgild/project/logs/current-batch.txt`
12 LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
13
14 echo "Starting daemons" >> $LOGFILE
15
16 start-all.sh
17 start-hbase.sh
18 mr-jobhistory-daemon.sh start historyserver
```

BIG DATA PROJECT

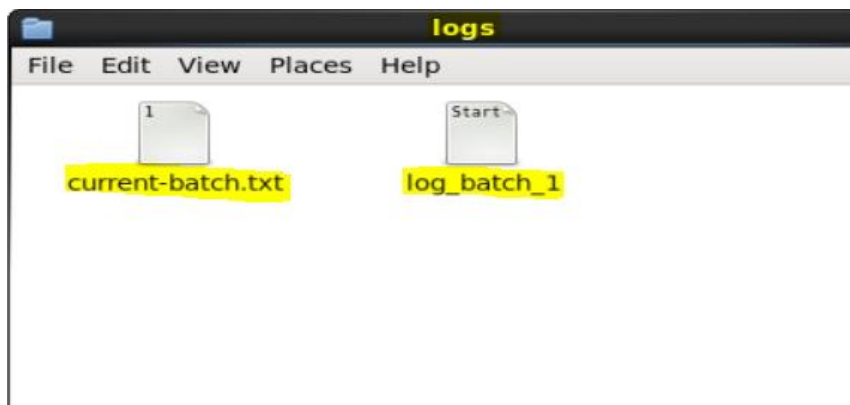
2. Starting all daemons.

The batch file script **start-daemons.sh** will start all the hadoop daemons and Hbase daemons as shown in the below screen shot.

```
[acadgild@localhost music]$ sh start-daemons.sh
After chmod
After batchid--> 1
This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh
19/01/20 19:08:43 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Starting namenodes on [localhost]
localhost: starting namenode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-namenode-localhost.localdomain.out
localhost: starting datanode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-datanode-localhost.localdomain.out
Starting secondary namenodes [0.0.0.0]
0.0.0.0: starting secondarynamenode, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/hadoop-acadgild-secondarynamenode-localhost.localdomain.out
19/01/20 19:09:31 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
starting yarn daemons
starting resourcemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/yarn-acadgild-resourcemanager-localhost.localdomain.out
localhost: starting nodemanager, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/yarn-acadgild-nodemanager-localhost.localdomain.out
localhost: starting zookeeper, logging to /home/acadgild/install/hbase/hbase-1.2.6/logs/hbase-acadgild-zookeeper-localhost.localdomain.out
starting master, logging to /home/acadgild/install/hbase/hbase-1.2.6/logs/hbase-acadgild-master-localhost.localdomain.out
starting regionserver, logging to /home/acadgild/install/hbase/hbase-1.2.6/logs/hbase-acadgild-1-regionserver-localhost.localdomain.out
starting historyserver, logging to /home/acadgild/install/hadoop/hadoop-2.6.5/logs/mapred-acadgild-historyserver-localhost.localdomain.out
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost music]$
```

```
[acadgild@localhost music]$ jps
4240 HQuorumPeer
3473 SecondaryNameNode
3250 DataNode
3622 ResourceManager
4887 Jps
4475 HRegionServer
4574 JobHistoryServer
3726 NodeManager
3151 NameNode
4319 HMaster
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost music]$
```

4. The **start-daemon.sh** script will check whether the current-batch.txt file is available in the logs folder or not. If not it will create the file and dump value '1' in that file and create LOGFILE with the current **batchid**.



SECTION – 4 – Data Ingestion, Formatting, Enrichment and Filtering

4.1 Stage – 1 – Data Ingestion

By using the “**populate-lookup.sh**” script we will create lookup tables in **Hbase**. These tables have to be used in,

- Data Formatting
- Data Enrichment
- Analysis Stage

Lookup Tables –

Table Name	Description
Station_Geo_Map	Contains mapping of a geo_cd with station_id
Subscribed_Users	Contains user_id, subscription_start_date and subscription_end_date. Contains details only for subscribed users
Song_Artist_Map	Contains mapping of song_id with artist_id alongwith royalty associated with each play of the song
User_Artist_Map	Contains an array of artist_id(s) followed by a user_id

The “**populate-lookup.sh**” shell script creates the above 4 lookup tables in the Hbase and populate the data into the lookup tables from the dataset files.

In the below screen shots, we can see the create-lookup.sh script and the following screen shots shows the table creation and population of the data in the Hbase. Also, the values loaded into the Hbase Tables are also shown, please see the below screen shots.


```

1  #!/bin/bash
2  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
3  LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
4  echo "Creating LookUp Tables" >> $LOGFILE
5  echo "create 'station-geo-map', 'geo'" | hbase shell
6  echo "create 'subscribed-users', 'subscn'" | hbase shell
7  echo "create 'song-artist-map', 'artist'" | hbase shell
8  echo "Populating LookUp Tables" >> $LOGFILE
9
10 file="/home/acadgild/project/lookupfiles/stn-geocd.txt"
11 while IFS= read -r line
12 do
13   stnid=`echo $line | cut -d',' -f1`
14   geocd=`echo $line | cut -d',' -f2`
15   echo "put 'station-geo-map', '$stnid', 'geo:geo_cd', '$geocd'" | hbase shell
16 done <"$file"
17
18 file="/home/acadgild/project/lookupfiles/song-artist.txt"
19 while IFS= read -r line
20 do
21   songid=`echo $line | cut -d',' -f1`
22   artistid=`echo $line | cut -d',' -f2`
23   echo "put 'song-artist-map', '$songid', 'artist:artistid', '$artistid'" | hbase shell
24 done <"$file"
25
26 file="/home/acadgild/project/lookupfiles/user-subscn.txt"
27 while IFS= read -r line
28 do
29   userid=`echo $line | cut -d',' -f1`
30   startdt=`echo $line | cut -d',' -f2`
31   enddt=`echo $line | cut -d',' -f3`
32   echo "put 'subscribed-users', '$userid', 'subscn:startdt', '$startdt'" | hbase shell
33   echo "put 'subscribed-users', '$userid', 'subscn:enddt', '$enddt'" | hbase shell
34 done <"$file"
35
36 hive -f /home/acadgild/project/scripts/user-artist.hql

```

Run Script: `./populate-lookup.sh`

```

[acadgild@localhost music]$ sh populate-lookup.sh
2019-01-20 19:26:06,756 WARN [main] util.NativeCodeLoader: Unable to load
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbase-1.2.6/
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.
s]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explan
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.6, rUnknown, Mon May 29 02:25:32 CDT 2017

create 'station-geo-map', 'geo'
0 row(s) in 3.4310 seconds

Hbase::Table - station-geo-map
2019-01-20 19:26:34,509 WARN [main] util.NativeCodeLoader: Unable to load
SLF4J: Class path contains multiple SLF4J bindings.

```

```

Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.6, rUnknown, Mon May 29 02:25:32 CDT 2017

create 'subscribed-users', 'subscn'
0 row(s) in 2.3100 seconds

Hbase::Table - subscribed-users
2019-01-20 19:26:59,822 WARN [main] util.NativeCodeLoader: Unable
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbas
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/had
s]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.6, rUnknown, Mon May 29 02:25:32 CDT 2017

create 'song-artist-map', 'artist'
0 row(s) in 3.2640 seconds

Hbase::Table - song-artist-map
2019-01-20 19:27:28,026 WARN [main] util.NativeCodeLoader: Unable
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hbase/hbas
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/had

```

We can see the lookup tables created using the *"populate-lookup.sh"* in the below screen shot, Lookup Tables in the hbase shell,

```

hbase(main):006:0> list
TABLE
song-artist-map
station-geo-map
subscribed-users
3 row(s) in 0.0110 seconds

=> ["song-artist-map", "station-geo-map", "subscribed-users"]
hbase(main):007:0>

```

Values loaded in lookup tables –

song-artist-map

```

hbase(main):007:0> scan "song-artist-map"
ROW COLUMN+CELL
S200 column=artist:artistid, timestamp=1547993035440, value=A300
S201 column=artist:artistid, timestamp=1547993063458, value=A301
S202 column=artist:artistid, timestamp=1547993091325, value=A302
S203 column=artist:artistid, timestamp=1547993118883, value=A303
S204 column=artist:artistid, timestamp=1547993145142, value=A304
S205 column=artist:artistid, timestamp=1547993170719, value=A301
S206 column=artist:artistid, timestamp=1547993196906, value=A302
S207 column=artist:artistid, timestamp=1547993221438, value=A303
S208 column=artist:artistid, timestamp=1547993246506, value=A304
S209 column=artist:artistid, timestamp=1547993269727, value=A305
10 row(s) in 0.3480 seconds

```

station-geo-map

```
hbase(main):008:0> scan "station-geo-map"
ROW                                COLUMN+CELL
ST400                             column=geo:geo_cd, timestamp=1547992654088, value=A
ST401                             column=geo:geo_cd, timestamp=1547992680289, value=AU
ST402                             column=geo:geo_cd, timestamp=1547992705729, value=AP
ST403                             column=geo:geo_cd, timestamp=1547992730660, value=J
ST404                             column=geo:geo_cd, timestamp=1547992756282, value=E
ST405                             column=geo:geo_cd, timestamp=1547992780470, value=A
ST406                             column=geo:geo_cd, timestamp=1547992805348, value=AU
ST407                             column=geo:geo_cd, timestamp=1547992829650, value=AP
ST408                             column=geo:geo_cd, timestamp=1547992854680, value=E
ST409                             column=geo:geo_cd, timestamp=1547992881341, value=E
ST410                             column=geo:geo_cd, timestamp=1547992906182, value=A
ST411                             column=geo:geo_cd, timestamp=1547992931229, value=A
ST412                             column=geo:geo_cd, timestamp=1547992956437, value=AP
ST413                             column=geo:geo_cd, timestamp=1547992982974, value=J
ST414                             column=geo:geo_cd, timestamp=1547993008451, value=E
15 row(s) in 0.2040 seconds
```

subscribed-users

```
hbase(main):009:0> scan "subscribed-users"
ROW                                COLUMN+CELL
U100                             column=subscn:enddt, timestamp=1547993317576, value=1465130523
U100                             column=subscn:startdt, timestamp=1547993294051, value=1465230523
U101                             column=subscn:enddt, timestamp=1547993365493, value=1475130523
U101                             column=subscn:startdt, timestamp=1547993342328, value=1465230523
U102                             column=subscn:enddt, timestamp=1547993413586, value=1475130523
U102                             column=subscn:startdt, timestamp=1547993389311, value=1465230523
U103                             column=subscn:enddt, timestamp=1547993461985, value=1475130523
U103                             column=subscn:startdt, timestamp=1547993437937, value=1465230523
U104                             column=subscn:enddt, timestamp=1547993508761, value=1475130523
U104                             column=subscn:startdt, timestamp=1547993485687, value=1465230523
U105                             column=subscn:enddt, timestamp=1547993555908, value=1475130523
U105                             column=subscn:startdt, timestamp=1547993532535, value=1465230523
U106                             column=subscn:enddt, timestamp=1547993602574, value=1485130523
U106                             column=subscn:startdt, timestamp=1547993578778, value=1465230523
U107                             column=subscn:enddt, timestamp=1547993651828, value=1455130523
U107                             column=subscn:startdt, timestamp=1547993626954, value=1465230523
U108                             column=subscn:enddt, timestamp=1547993698922, value=1465230623
U108                             column=subscn:startdt, timestamp=1547993675166, value=1465230523
U109                             column=subscn:enddt, timestamp=1547993745677, value=1475130523
U109                             column=subscn:startdt, timestamp=1547993722471, value=1465230523
U110                             column=subscn:enddt, timestamp=1547993793080, value=1475130523
U110                             column=subscn:startdt, timestamp=1547993769905, value=1465230523
U111                             column=subscn:enddt, timestamp=1547993840227, value=1475130523
U111                             column=subscn:startdt, timestamp=1547993816211, value=1465230523
U112                             column=subscn:enddt, timestamp=1547993888721, value=1475130523
U112                             column=subscn:startdt, timestamp=1547993864039, value=1465230523
U113                             column=subscn:enddt, timestamp=1547993937535, value=1485130523
U113                             column=subscn:startdt, timestamp=1547993913032, value=1465230523
U114                             column=subscn:enddt, timestamp=1547993985752, value=1468130523
U114                             column=subscn:startdt, timestamp=1547993961771, value=1465230523
15 row(s) in 0.2710 seconds
```

We have successfully created the lookup tables in the Hbase.

The populate-lookup.sh also creates a lookup table “**users_artists**” in the HIVE, loading the data from the **user-artist.txt**, the below screen shot shows that the table has been created in the HIVE.

```

hive> show databases;
OK
default
emp
olympic
project
Time taken: 0.109 seconds, Fetched: 4 row(s)
hive> use project;
OK
Time taken: 0.098 seconds
hive> show tables;
OK
users_artists
Time taken: 0.149 seconds, Fetched: 1 row(s)
hive> select * from users_artists;
OK
U100      ["A300","A301","A302"]
U101      ["A301","A302"]
U102      ["A302"]
U103      ["A303","A301","A302"]
U104      ["A304","A301"]
U105      ["A305","A301","A302"]
U106      ["A301","A302"]
U107      ["A302"]
U108      ["A300","A303","A304"]
U109      ["A301","A303"]
U110      ["A302","A301"]
U111      ["A303","A301"]
U112      ["A304","A301"]
U113      ["A305","A302"]
U114      ["A300","A301","A302"]
Time taken: 0.608 seconds, Fetched: 15 row(s)

```

Now we need to link these lookup tables in hive using the Hbase Storage Handler.

With the help of “data_enrichment_filtering_schema.sh” file we will create hive tables on the top of Hbase tables using “create_hive_hbase_lookup.hql”.

Creating Hive tables on top of HBASE-

In this section with the help of Hbase storage handler & SerDe properties we are creating the hive external tables by matching the columns of Hbase tables to hive tables.

Run the script: ./data_enrichment_filtering_schema.sh

The script will run the “create_hive_hbase_lookup.hql” which will create the HIVE external tables with the help of **Hbase storage handler & SerDe properties**. The hive external tables will match the columns of **Hbase** tables to **HIVE** tables.

```

1  #!/bin/bash
2
3  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
4  LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
5
6  echo "Creating hive tables on top of hbase tables for data enrichment and filtering..." >> $LOGFILE
7
8  hive -f /home/acadgild/project/scripts/create_hive_hbase_lookup.hql

```

create_hive_hbase_lookup.hql

```

1  USE project;
2  create external table if not exists station_geo_map
3  (
4  station_id String,
5  geo_cd string
6  )
7  STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
8  with serdeproperties
9  ("hbase.columns.mapping"=":key,geo:geo_cd")
10 tblproperties("hbase.table.name"="station-geo-map");
11
12 create external table if not exists subscribed_users
13 (
14 user_id STRING,
15 subscn_start_dt STRING,
16 subscn_end_dt STRING
17 )
18 STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
19 with serdeproperties
20 ("hbase.columns.mapping"=":key,subscn:startdt,subscn:enddt")
21 tblproperties("hbase.table.name"="subscribed-users");
22
23 create external table if not exists song_artist_map
24 (
25 song_id STRING,
26 artist_id STRING
27 )
28 STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
29 with serdeproperties
30 ("hbase.columns.mapping"=":key,artist:artistid")
31 tblproperties("hbase.table.name"="song-artist-map");

```

The below screenshot we can see tables getting created in hive by running the *"data_enrichement_filtering_schema.sh file"*

```

hive> use project;
OK
Time taken: 12.272 seconds
hive> show tables;
OK
song_artist_map
station_geo_map
subscribed_users
users_artists
Time taken: 0.619 seconds, Fetched: 4 row(s)
hive> █

```

*Select * from song_artist_map*

```
hive> select * from song_artist_map;
OK
S200      A300
S201      A301
S202      A302
S203      A303
S204      A304
S205      A301
S206      A302
S207      A303
S208      A304
S209      A305
Time taken: 0.791 seconds, Fetched: 10 row(s)
```

*Select * from station_geo_map;*

```
hive> select * from station_geo_map;
OK
ST400     A
ST401     AU
ST402     AP
ST403     J
ST404     E
ST405     A
ST406     AU
ST407     AP
ST408     E
ST409     E
ST410     A
ST411     A
ST412     AP
ST413     J
ST414     E
Time taken: 0.713 seconds, Fetched: 15 row(s)
```

*Select * from subscribed_users;*

```
hive> select * from subscribed_users;
OK
U100      1465230523      1465130523
U101      1465230523      1475130523
U102      1465230523      1475130523
U103      1465230523      1475130523
U104      1465230523      1475130523
U105      1465230523      1475130523
U106      1465230523      1485130523
U107      1465230523      1455130523
U108      1465230523      1465230623
U109      1465230523      1475130523
U110      1465230523      1475130523
U111      1465230523      1475130523
U112      1465230523      1475130523
U113      1465230523      1485130523
U114      1465230523      1468130523
Time taken: 0.768 seconds, Fetched: 15 row(s)
```


4.2 Stage-2 – Data Formatting

In this stage we are merging the data coming from both **web** applications and **mobile** applications and create a common table for analyzing purpose and create partitioned data based on **batchid**, since we are running this script for every 3 hours.

Run the script: `./dataformatting.sh`

```

1  #!/bin/bash
2
3  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
4  LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
5
6  echo "Placing data files from local to HDFS..." >> $LOGFILE
7
8  hadoop fs -rm -r /user/acadgild/project/batch${batchid}/web/
9  hadoop fs -rm -r /user/acadgild/project/batch${batchid}/formattedweb/
10 hadoop fs -rm -r /user/acadgild/project/batch${batchid}/mob/
11
12 hadoop fs -mkdir -p /user/acadgild/project/batch${batchid}/web/
13 hadoop fs -mkdir -p /user/acadgild/project/batch${batchid}/mob/
14
15 hadoop fs -put /home/acadgild/project/data/web/* /user/acadgild/project/batch${batchid}/web/
16 hadoop fs -put /home/acadgild/project/data/mob/* /user/acadgild/project/batch${batchid}/mob/
17
18 echo "Running pig script for data formatting..." >> $LOGFILE
19
20 pig -param batchid=${batchid} /home/acadgild/project/scripts/dataformatting.pig
21
22 echo "Running hive script for formatted data load..." >> $LOGFILE
23
24 hive -hiveconf batchid=${batchid} -f /home/acadgild/project/scripts/formatted_hive_load.hql

```

formatted_hive_load.hql

```

1  USE project;
2
3  CREATE TABLE IF NOT EXISTS formatted_input
4  (
5    user_id STRING,
6    song_id STRING,
7    artist_id STRING,
8    time_stamp STRING,
9    start_ts STRING,
10   end_ts STRING,
11   geo_cd STRING,
12   station_id STRING,
13   song_end_type INT,
14   liked INT,
15   disliked INT
16  )
17  ROW FORMAT DELIMITED
18  FIELDS TERMINATED BY ','
19  PARTITIONED BY (batchid INT);
20  LOAD DATA INPATH '/user/acadgild/project/batch${hiveconf:batchid}/formattedweb/'
21  INTO TABLE formatted_input PARTITION (batchid=${hiveconf:batchid});
22  LOAD DATA INPATH '/user/acadgild/project/batch${hiveconf:batchid}/mob/'
23  INTO TABLE formatted_input PARTITION (batchid=${hiveconf:batchid});

```

In the below screenshot we can see the data both the scripts in action, first pig script will parse the data and then hive script will load the data into hive terminal successfully.

BIG DATA PROJECT

Pig script completion -

```
2019-01-20 21:09:49,867 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2019-01-20 21:09:49,879 [main] INFO org.apache.pig.tools.pigstats.mapreduce.SimplePigStats - Script Statistics:

HadoopVersion PigVersion   UserId  StartedAt       FinishedAt       Features
2.6.5         0.16.0   acadgild 2019-01-20 21:08:00 2019-01-20 21:09:49 UNKNOWN

Success!

Job Stats (time in seconds):
JobId  Maps  Reduces MaxMapTime   MinMapTime   AvgMapTime   MedianMapTime   MaxReduceTime   MinReduceTime   AvgReduceTime
s      Feature Outputs
job_1547991585611_0001 1      0      35      35      35      35      0      0      0      0      A,B      MAP_ONLY      /user/
tedweb,

Input(s):
Successfully read 200 records (67348 bytes) from: "/user/acadgild/project/batch1/web"

Output(s):
Successfully stored 200 records (12338 bytes) in: "/user/acadgild/project/batch1/formattedweb"

Counters:
Total records written : 200
Total bytes written : 12338
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1547991585611_0001
```

Hive script successfully load the data into hive terminal –

```
Logging initialized using configuration in jar:file:/home/acadgild/install/hive/apache-hive
OK
Time taken: 18.1 seconds
OK
Time taken: 1.808 seconds
Loading data to table project.formatted_input partition (batchid=1)
OK
Time taken: 5.602 seconds
Loading data to table project.formatted_input partition (batchid=1)
OK
Time taken: 2.597 seconds
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost music]$
```

In the above screenshot we can see the **dataformatting.pig** along with the **formatted_hive_load.hql** executed successfully.

The output of **dataformatting.sh** script in HDFS folders:

```
[acadgild@localhost music]$ hadoop fs -ls /user/acadgild/project/batch1
19/01/23 03:55:44 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
Found 3 items
drwxr-xr-x - acadgild supergroup          0 2019-01-23 03:54 /user/acadgild/project/batch1/formattedweb
drwxr-xr-x - acadgild supergroup          0 2019-01-20 22:21 /user/acadgild/project/batch1/mob
drwxr-xr-x - acadgild supergroup          0 2019-01-23 03:52 /user/acadgild/project/batch1/web
[acadgild@localhost music]$ hadoop fs -ls /user/acadgild/project/batch1/formattedweb
19/01/23 03:56:19 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java
Found 2 items
-rw-r--r-- 1 acadgild supergroup          0 2019-01-23 03:54 /user/acadgild/project/batch1/formattedweb/_SUCCESS
-rw-r--r-- 1 acadgild supergroup     12338 2019-01-23 03:53 /user/acadgild/project/batch1/formattedweb/part-m-00000
[acadgild@localhost music]$
```



```
[acadgild@localhost music]$ hadoop fs -cat /user/acadgild/project/batch1/formattedweb/*
19/01/23 03:57:26 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform.
U104,S205,A304,1462863262,1462863262,1494297562,AU,ST410,1,0,0
U104,S200,A305,1494297562,1465490556,1465490556,AU,ST412,1,0,1
U111,S201,,1494297562,1465490556,1468094889,U,ST404,1,0,0
U104,S200,A300,1468094889,1465490556,1494297562,AP,ST409,1,1,1
U113,S203,A305,1468094889,1462863262,1468094889,U,ST414,3,0,0
,S204,A302,1468094889,1494297562,1494297562,A,ST400,3,1,0
U108,S209,A302,1468094889,1465490556,1494297562,AU,ST410,2,0,1
U111,S207,A302,1462863262,1468094889,1465490556,E,ST404,2,0,0
U102,S204,A303,1465490556,1465490556,1468094889,,ST411,1,1,1
U113,S206,A305,1468094889,1494297562,1462863262,U,ST409,1,1,0
U112,S204,A302,1494297562,1468094889,1468094889,AU,ST411,0,1,0
U115,S205,A305,1465490556,1494297562,1494297562,E,ST412,3,1,1
U105,S205,A300,1462863262,1462863262,1462863262,A,ST402,2,1,0
U115,S201,,1465490556,1462863262,1494297562,U,ST402,3,1,0
U113,S203,A302,1465490556,1462863262,1465490556,AU,ST413,3,1,1
U112,S205,A305,1494297562,1465490556,1494297562,U,ST407,0,1,1
U107,S210,A300,1465490556,1465490556,1462863262,AP,ST406,2,0,0
U103,S209,A301,1465490556,1494297562,1468094889,AU,ST411,2,1,0
U114,S209,A303,1494297562,1468094889,1465490556,A,ST405,0,0,1
U104,S210,A305,1468094889,1494297562,1465490556,AU,ST413,1,1,1
,S201,A300,1468094889,1468094889,1494297562,,ST409,1,0,0
U111,S203,A300,1468094889,1465490556,1465490556,AP,ST412,3,0,1
U108,S208,A301,1494297562,1465490556,1494297562,AP,ST411,1,0,0
U107,S202,A302,1465490556,1468094889,1465490556,E,ST409,3,1,0
U112,S204,,1494297562,1462863262,1462863262,AP,ST410,2,0,1
U106,S205,A301,1465490556,1494297562,1494297562,E,ST415,2,1,0
U113,S209,A302,1468094889,1468094889,1468094889,U,ST410,3,0,0
U108,S209,A301,1465490556,1462863262,1468094889,AU,ST415,0,0,1
U109,S210,A302,1465490556,1462863262,1465490556,AP,ST408,1,1,1
U101,S207,A303,1468094889,1462863262,1494297562,AU,ST402,3,1,1
U107,S209,A301,1465490556,1494297562,1462863262,A,ST402,2,0,0
U110,S206,A302,1468094889,1465490556,1465490556,U,ST409,2,1,0
```

The *formatted_input* table is created in HIVE after running **dataformatting.sh** script,

```
hive> use project;
OK
Time taken: 17.537 seconds
hive> show tables;
OK
formatted_input
song_artist_map
station_geo_map
subscribed_users
users_artists
Time taken: 0.858 seconds, Fetched: 5 row(s)
hive> select * from formatted_input;
OK
U114 S209 A303 1495130523 1465230523 1475130523 AP ST402 3 1 1 1
U110 S205 A301 1475130523 1465130523 1465130523 A ST401 1 0 1 1
U120 S200 1495130523 1475130523 1475130523 E ST401 2 0 0 1
U105 S208 A301 1475130523 1465130523 1475130523 E ST408 1 1 0 1
U107 S208 A301 1465130523 1485130523 1465130523 A ST412 0 0 0 1
,S200 A302 1495130523 1465130523 1465230523 U ST405 3 0 0 1
U102 S205 A304 1465230523 1475130523 1465230523 AU ST409 0 1 1 1
U101 S206 A305 1465130523 1475130523 1465230523 AP ST415 2 1 0 1
U108 S203 A303 1475130523 1465230523 1475130523 ST400 1 1 1 1
U119 S207 A305 1465230523 1485130523 1485130523 E ST409 2 0 0 1
U104 S206 A300 1475130523 1475130523 1465230523 A ST406 3 0 0 1
U115 S207 A302 1465130523 1465230523 1475130523 A ST409 2 1 1 1
U101 S201 A304 1495130523 1465230523 1475130523 E ST403 3 1 0 1
U102 S210 1465130523 1465230523 1475130523 U ST409 3 0 1 1
U114 S205 A302 1465130523 1475130523 1485130523 A ST415 3 0 0 1
U117 S210 A300 1495130523 1485130523 1485130523 U ST400 0 1 1 1
U117 S201 A301 1465130523 1465230523 1465230523 AP ST413 0 1 0 1
U112 S206 A302 1465230523 1475130523 1485130523 AU ST409 1 0 1 1
U119 S204 A300 1465230523 1465230523 1465230523 AP ST406 3 1 1 1
U115 S207 A300 1475130523 1475130523 1485130523 AU ST413 3 0 1 1
,S204 A301 1465130523 1485130523 1485130523 ST406 0 1 1 1
U104 S206 A303 1465130523 1465230523 1465230523 A ST401 1 1 0 1
U103 S210 A302 1475130523 1485130523 1485130523 U ST402 1 0 1 1
U108 S207 A304 1495130523 1475130523 1475130523 A ST415 1 0 0 1
U108 S209 1465230523 1465130523 1465130523 A ST403 1 1 0 1
```

- In the above screenshot we can see the formatted input data with some null values in **user_id**, **artist_id** and **geo_cd** columns which we will fill in stage-3 (Data Enrichment) based on rules of enrichment for **artist_id** and **geo_cd** only. We will get neglect **user_id** because they didn't mention anything about **user_id** for enrichment purpose.

- Data formatting phase is executed successfully by loading both **mobile** and **web** data and partitioned based on **batchid**.

4.3 Stage – 3 – Data Enrichment & Filtering

In this stage, we will enrich the data coming from **web** and **mobile** applications using the lookup table stored in **Hbase** and divide the records based on the enrichment rules into 'pass' and 'fail' records.

Rules for data enrichment -

- If any of like or dislike is **NULL** or **absent**, consider it as **0**.
- If fields like **Geo_cd** and **Artist_id** are NULL or absent, consult the lookup tables for fields **Station_id** and **Song_id** respectively to get the values of **Geo_cd** and **Artist_id**.
- If corresponding lookup entry is not found, consider that **record** to be **invalid**

So based on the enrichment rules we will fill the null **geo_cd** and **artist_id** values with the help of corresponding lookup values in **song-artist-map** and **station-geo-map** tables in **Hive-Hbase** tables.

data_enrichment.sh

```

1  #!/bin/bash
2
3  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
4  LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
5  VALIDDIR=/home/acadgild/project/processed_dir/valid/batch_${batchid}
6  INVALIDDIR=/home/acadgild/project/processed_dir/invalid/batch_${batchid}
7
8  echo "Running hive script for data enrichment and filtering..." >> $LOGFILE
9
10 hive -hiveconf batchid=${batchid} -f /home/acadgild/project/scripts/data_enrichment.hql
11
12 if [ ! -d "$VALIDDIR" ]
13 then
14 mkdir -p "$VALIDDIR"
15 fi
16
17 if [ ! -d "$INVALIDDIR" ]
18 then
19 mkdir -p "$INVALIDDIR"
20 fi
21
22 echo "Copying valid and invalid records in local file system..." >> $LOGFILE
23
24 hadoop fs -get /user/hive/warehouse/project.db/enriched_data/batchid=${batchid}/status=pass/* $VALIDDIR
25 hadoop fs -get /user/hive/warehouse/project.db/enriched_data/batchid=${batchid}/status=fail/* $INVALIDDIR
26
27 echo "Deleting older valid and invalid records from local file system..." >> $LOGFILE
28
29 find /home/acadgild/project/processed_dir/ -mtime +7 -exec rm {} \;
```

data_enrichment.hql

```

1  SET hive.auto.convert.join=false;
2  SET hive.exec.dynamic.partition.mode=nonstrict;
3  USE project;
4  CREATE TABLE IF NOT EXISTS enriched_data
5  (
6  User_id STRING,
7  Song_id STRING,
8  Artist_id STRING,
9  time_stamp STRING,
10 Start_ts STRING,
11 End_ts STRING,
12 Geo_cd STRING,
13 Station_id STRING,
14 Song_end_type INT,
15 liked INT,
16 disliked INT
17 )
18 PARTITIONED BY
19 (batchid INT,
20 status STRING)
21 STORED AS ORC;
22
23 INSERT OVERWRITE TABLE enriched_data
24 PARTITION (batchid, status)
25 SELECT
26 i.user_id,
27 i.song_id,
28 sa.artist_id,
29 i.time_stamp,
30 i.start_ts,
31 i.end_ts,
32 sg.geo_cd,
33 i.station_id,
34 IF (i.song_end_type IS NULL, 3, i.song_end_type) AS song_end_type,
35 IF (i.liked IS NULL, 0, i.liked) AS liked,
36 IF (i.disliked IS NULL, 0, i.disliked) AS disliked,
37 i.batchid,
38 IF((i.liked=1 AND i.disliked=1)
39 OR i.user_id IS NULL
40 OR i.song_id IS NULL
41 OR i.time_stamp IS NULL
42 OR i.start_ts IS NULL
43 OR i.end_ts IS NULL
44 OR i.geo_cd IS NULL
45 OR i.user_id=''
46 OR i.song_id=''
47 OR i.time_stamp=''
48 OR i.start_ts=''
49 OR i.end_ts=''
50 OR i.geo_cd='')
51 OR sg.geo_cd IS NULL
52 OR sg.geo_cd=''
53 OR sa.artist_id IS NULL
54 OR sa.artist_id='', 'fail', 'pass') AS status
55 FROM formatted_input i
56 LEFT OUTER JOIN station_geo_map sg ON i.station_id = sg.station_id
57 LEFT OUTER JOIN song_artist_map sa ON i.song_id = sa.song_id
58 WHERE i.batchid=${hiveconf:batchid};

```

BIG DATA PROJECT

Run Script: `./data_enrichment.sh`

```
[acadgild@localhost music]$ sh data_enrichment.sh
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12.jar]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j.properties
OK
Time taken: 17.226 seconds
OK
Time taken: 1.827 seconds
No Stats for project@formatted_input, Columns: start_ts, song_id, time_stamp, user_id, end_ts, station_id, geo_cd,
No Stats for project@station_geo_map, Columns: station_id, geo_cd
No Stats for project@song_artist_map, Columns: song_id, artist_id
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff
Hive 1.X releases.
Query ID = acadgild_20190123042026_6c6b0de1-1d8f-453e-b177-33f54e28d82d
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1548195624342_0003, Tracking URL = http://localhost:8088/proxy/application_1548195624342_0003/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1548195624342_0003
```

```
Hadoop job information for Stage-2: number of mappers: 2; number of reducers: 1
2019-01-23 04:24:54,417 Stage-2 map = 0%, reduce = 0%
2019-01-23 04:25:40,547 Stage-2 map = 50%, reduce = 0%, Cumulative CPU 8.76 sec
2019-01-23 04:25:45,655 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 20.51 sec
2019-01-23 04:26:19,540 Stage-2 map = 100%, reduce = 67%, Cumulative CPU 26.07 sec
2019-01-23 04:26:25,944 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 32.17 sec
MapReduce Total cumulative CPU time: 32 seconds 170 msec
Ended Job = job_1548195624342_0004
Loading data to table project.enriched_data partition (batchid=null, status=null)

Loaded : 2/2 partitions.
Time taken to load dynamic partitions: 4.42 seconds
Time taken for adding to write entity : 0.017 seconds
MapReduce Jobs Launched:
Stage-Stage-1: Map: 3 Reduce: 1 Cumulative CPU: 38.62 sec HDFS Read: 71749 HDFS Write: 30063 SUCCESS
Stage-Stage-2: Map: 2 Reduce: 1 Cumulative CPU: 32.17 sec HDFS Read: 51146 HDFS Write: 5176 SUCCESS
Total MapReduce CPU Time Spent: 1 minutes 10 seconds 790 msec
OK
Time taken: 370.938 seconds
19/01/23 04:26:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
19/01/23 04:26:52 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost music]$
```

At the end script will automatically divide the records based on status **pass** & **fail** and dump the result into **processed_dir** folder with valid and invalid folders.

```
[acadgild@localhost processed_dir]$ ls -l
total 8
drwxrwxr-x. 3 acadgild acadgild 4096 Jan 23 04:11 invalid
drwxrwxr-x. 3 acadgild acadgild 4096 Jan 23 04:11 valid
[acadgild@localhost processed_dir]$ ls -l valid/batch_1/
total 4
-rw-r--r--. 1 acadgild acadgild 2590 Jan 23 04:26 000000_0
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost processed_dir]$ ls -l invalid/batch_1/
total 4
-rw-r--r--. 1 acadgild acadgild 2415 Jan 23 04:26 000000_0
[acadgild@localhost processed_dir]$
```


BIG DATA PROJECT

Now we can check whether the data properly loaded in the HIVE.

```
hive> show tables;
OK
enriched_data
formatted_input
song_artist_map
station_geo_map
subscribed_users
users_artists
Time taken: 8.336 seconds, Fetched: 6 row(s)
hive> █
```

In the below screenshot we have data for **enriched_data** table where we filled the null values of **artist_id** and **geo_cd** of formatted input with the help of lookup tables.

*select * from enriched_data;*

```
hive> select * from enriched_data;
OK
U108 S200 A300 1465490556 1462863262 1462863262 A ST400 1 1 1 1 fail
S200 A300 1495130523 1465130523 1465230523 A ST405 3 0 0 1 fail
U115 S200 A300 1465130523 1465230523 1485130523 E ST404 3 1 1 1 fail
U120 S200 A300 1465230523 1475130523 1485130523 AP ST407 2 1 1 1 fail
U105 S200 A300 1468094889 1494297562 1465490556 AP ST402 0 1 1 1 fail
U116 S200 A300 1495130523 1475130523 1485130523 J ST403 2 1 1 1 fail
S200 A300 1462863262 1462863262 1494297562 AP ST402 2 1 0 1 fail
U118 S200 A300 1495130523 1485130523 1465130523 AP ST402 0 1 1 1 fail
U104 S200 A300 1468094889 1465490556 1494297562 E ST409 1 1 1 1 fail
U103 S200 A300 1465130523 1475130523 1465230523 E ST409 2 1 1 1 fail
U109 S200 A300 1495130523 1465230523 1465130523 E ST409 3 1 1 1 fail
U101 S200 A300 1465230523 1475130523 1465130523 AU ST401 2 1 1 1 fail
U106 S200 A300 1465230523 1465230523 1465230523 AU ST401 0 1 1 1 fail
U107 S200 A300 1475130523 1475130523 1485130523 AU ST401 2 1 1 1 fail
U114 S200 A300 1468094889 1465490556 1494297562 A ST400 2 1 1 1 fail
U118 S200 A300 1465490556 1462863262 1494297562 A ST400 1 1 1 1 fail
S201 A301 1494297562 1465490556 1494297562 J ST413 1 1 0 1 fail
S201 A301 1465230523 1465230523 1465230523 A ST405 0 1 0 1 fail
S201 A301 1494297562 1468094889 1462863262 E ST414 3 1 0 1 fail
U112 S201 A301 1495130523 1465130523 1465230523 E ST409 0 1 1 1 fail
U114 S201 A301 1465130523 1475130523 1485130523 A ST410 0 1 1 1 fail
U111 S201 A301 1465130523 1475130523 1465130523 E ST408 0 0 0 1 fail
U119 S201 A301 1465490556 1465490556 1494297562 AP ST402 2 1 1 1 fail
U116 S201 A301 1495130523 1485130523 1475130523 E ST408 1 1 1 1 fail
U106 S201 A301 1494297562 1465490556 1494297562 E ST408 1 1 1 1 fail
S201 A301 1468094889 1468094889 1494297562 E ST409 1 0 0 1 fail
U119 S202 A302 1462863262 1465490556 1494297562 E ST414 2 1 1 1 fail
U102 S202 A302 1462863262 1468094889 1494297562 NULL ST415 3 0 1 1 fail
U115 S202 A302 1475130523 1465230523 1465130523 A ST400 2 1 1 1 fail
U101 S202 A302 1465490556 1462863262 1462863262 A ST411 3 1 0 1 fail
U114 S202 A302 1495130523 1485130523 1465230523 AU ST401 1 1 1 1 fail
U113 S202 A302 1465230523 1465230523 1465230523 NULL ST415 1 1 1 1 fail
U100 S202 A302 1494297562 1462863262 1462863262 E ST409 1 1 1 1 fail
S202 A302 1475130523 1465230523 1465230523 AP ST407 0 1 1 1 fail
U102 S202 A302 1468094889 1465490556 1494297562 E ST414 3 1 1 1 fail
U109 S202 A302 1465490556 1465490556 1465490556 AU ST406 1 0 1 1 fail
U116 S208 A304 1462863262 1462863262 1465490556 E ST409 2 0 1 1 pass
U110 S208 A304 1462863262 1465490556 1494297562 E ST409 2 0 0 1 pass
U105 S208 A304 1465130523 1475130523 1465230523 AP ST402 1 0 1 1 pass
U103 S208 A304 1465230523 1475130523 1485130523 AP ST402 0 0 1 1 pass
U112 S208 A304 1465230523 1465230523 1465230523 A ST410 1 0 0 1 pass
U115 S208 A304 1495130523 1475130523 1465230523 AU ST401 1 0 1 1 pass
U116 S208 A304 1494297562 1468094889 1494297562 A ST410 2 1 0 1 pass
U108 S208 A304 1494297562 1465490556 1494297562 A ST411 1 0 0 1 pass
U104 S208 A304 1465230523 1475130523 1465130523 A ST411 1 0 1 1 pass
U113 S208 A304 1495130523 1475130523 1475130523 AP ST412 2 0 1 1 pass
U115 S208 A304 1465230523 1475130523 1485130523 AP ST412 0 0 1 1 pass
U107 S208 A304 1465130523 1485130523 1465130523 AP ST412 0 0 1 1 pass
U111 S208 A304 1462863262 1462863262 1462863262 A ST400 1 0 1 1 pass
U114 S208 A304 1468094889 1465490556 1462863262 E ST414 2 0 0 1 pass
U107 S208 A304 1465490556 1468094889 1494297562 E ST414 1 0 0 1 pass
U116 S208 A304 1465490556 1462863262 1494297562 E ST404 2 0 1 1 pass
U118 S209 A305 1465490556 1468094889 1468094889 E ST409 3 0 1 1 pass
U105 S209 A305 1495130523 1485130523 1485130523 J ST413 2 1 0 1 pass
U105 S209 A305 1465130523 1465130523 1485130523 AP ST402 3 1 0 1 pass
U105 S209 A305 1475130523 1465230523 1485130523 AP ST402 0 0 1 1 pass
U115 S209 A305 1465130523 1465130523 1485130523 E ST409 2 0 1 1 pass
U120 S209 A305 1494297562 1465490556 1494297562 E ST414 2 0 0 1 pass
U113 S209 A305 1468094889 1468094889 1468094889 A ST410 3 0 0 1 pass
U113 S209 A305 1495130523 1465130523 1485130523 E ST404 0 0 1 1 pass
U108 S209 A305 1468094889 1465490556 1494297562 A ST410 2 0 1 1 pass
U104 S209 A305 1494297562 1465490556 1465490556 A ST411 3 0 0 1 pass
U103 S209 A305 1465490556 1494297562 1468094889 A ST411 2 1 0 1 pass
U108 S209 A305 1465230523 1465130523 1465130523 J ST403 1 1 0 1 pass
U111 S209 A305 1475130523 1465130523 1465130523 AP ST412 3 0 1 1 pass
U112 S209 A305 1495130523 1485130523 1475130523 A ST400 3 0 0 1 pass
U107 S209 A305 1462863262 1468094889 1468094889 AP ST412 2 1 0 1 pass
U114 S209 A305 1494297562 1468094889 1465490556 A ST405 0 0 1 1 pass
U100 S209 A305 1468094889 1462863262 1494297562 J ST413 1 0 1 1 pass
U107 S209 A305 1465490556 1494297562 1462863262 AP ST402 2 0 0 1 pass
U104 S209 A305 1475130523 1465230523 1465230523 E ST408 3 1 0 1 pass
U119 S209 A305 1465230523 1485130523 1465130523 AP ST402 1 0 0 1 pass
Time taken: 5.406 seconds, Fetched: 400 row(s)
hive> █
```

By applying the mentioned rules, we have successfully accomplished Data enrichment and Filtering stage.

4.4 Stage – 4 – Data Analysis

In this stage we will do analysis on enriched data.

data_analysis.sh

```
1  #!/bin/bash
2
3  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
4  LOGFILE=/home/acadgild/project/logs/log_batch_${batchid}
5
6  echo "Running hive script for data analysis..." >> $LOGFILE
7
8  hive -hiveconf batchid=$batchid -f /home/acadgild/project/scripts/data_analysis.hql
9
10 sh /home/acadgild/project/scripts/data_export.sh
11
12 echo "Incrementing batchid..." >> $LOGFILE
13
14 batchid=`expr $batchid + 1`
15 echo -n $batchid > /home/acadgild/project/logs/current-batch.txt
```

data_analysis.hql

```
1  SET hive.auto.convert.join=false;
2  USE project;
3
4  CREATE TABLE IF NOT EXISTS top_10_stations
5  (
6  station_id STRING,
7  total_distinct_songs_played INT,
8  distinct_user_count INT
9  )
10 PARTITIONED BY (batchid INT)
11 ROW FORMAT DELIMITED
12 FIELDS TERMINATED BY ','
13 STORED AS TEXTFILE;
14
15 INSERT OVERWRITE TABLE top_10_stations
16 PARTITION(batchid=${hiveconf:batchid})
17 SELECT
18 station_id,
19 COUNT(DISTINCT song_id) AS total_distinct_songs_played,
20 COUNT(DISTINCT user_id) AS distinct_user_count
21 FROM enriched_data
22 WHERE status='pass'
23 AND batchid=${hiveconf:batchid}
24 AND liked=1
25 GROUP BY station_id
26 ORDER BY total_distinct_songs_played DESC
27 LIMIT 10;
```

```

30 CREATE TABLE IF NOT EXISTS users_behaviour
31 (
32   user_type STRING,
33   duration INT
34 )
35 PARTITIONED BY (batchid INT)
36 ROW FORMAT DELIMITED
37 FIELDS TERMINATED BY ','
38 STORED AS TEXTFILE;
39
40 INSERT OVERWRITE TABLE users_behaviour
41 PARTITION(batchid=${hiveconf:batchid})
42 SELECT
43   CASE WHEN (su.user_id IS NULL OR CAST(ed.timestamp AS DECIMAL(20,0)) > CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'UNSUBSCRIBED'
44   WHEN (su.user_id IS NOT NULL AND CAST(ed.timestamp AS DECIMAL(20,0)) <= CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'SUBSCRIBED'
45   END AS user_type,
46   SUM(ABS(CAST(ed.end_ts AS DECIMAL(20,0))-CAST(ed.start_ts AS DECIMAL(20,0)))) AS duration
47 FROM enriched_data ed
48 LEFT OUTER JOIN subscribed_users su
49 ON ed.user_id=su.user_id
50 WHERE ed.status='pass'
51 AND ed.batchid=${hiveconf:batchid}
52 GROUP BY CASE WHEN (su.user_id IS NULL OR CAST(ed.timestamp AS DECIMAL(20,0)) > CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'UNSUBSCRIBED'
53 WHEN (su.user_id IS NOT NULL AND CAST(ed.timestamp AS DECIMAL(20,0)) <= CAST(su.subscn_end_dt AS DECIMAL(20,0))) THEN 'SUBSCRIBED' END;
54

```

```

56 CREATE TABLE IF NOT EXISTS connected_artists
57 (
58   artist_id STRING,
59   user_count INT
60 )
61 PARTITIONED BY (batchid INT)
62 ROW FORMAT DELIMITED
63 FIELDS TERMINATED BY ','
64 STORED AS TEXTFILE;
65
66 INSERT OVERWRITE TABLE connected_artists
67 PARTITION(batchid=${hiveconf:batchid})
68 SELECT
69   ua.artist_id,
70   COUNT(DISTINCT ua.user_id) AS user_count
71 FROM
72   (
73     SELECT user_id, artist_id FROM users_artists
74     LATERAL VIEW explode(artists_array) artists AS artist_id
75   ) ua
76 INNER JOIN
77   (
78     SELECT artist_id, song_id, user_id
79     FROM enriched_data
80     WHERE status='pass'
81     AND batchid=${hiveconf:batchid}
82   ) ed
83 ON ua.artist_id=ed.artist_id
84 AND ua.user_id=ed.user_id
85 GROUP BY ua.artist_id
86 ORDER BY user_count DESC
87 LIMIT 10;

```

BIG DATA PROJECT

```
90 CREATE TABLE IF NOT EXISTS top_10_royalty_songs
91 (
92 song_id STRING,
93 duration INT
94 )
95 PARTITIONED BY (batchid INT)
96 ROW FORMAT DELIMITED
97 FIELDS TERMINATED BY ','
98 STORED AS TEXTFILE;
99
100 INSERT OVERWRITE TABLE top_10_royalty_songs
101 PARTITION(batchid=${hiveconf:batchid})
102 SELECT song_id,
103 SUM(ABS(CAST(end_ts AS DECIMAL(20,0))-CAST(start_ts AS DECIMAL(20,0)))) AS duration
104 FROM enriched_data
105 WHERE status='pass'
106 AND batchid=${hiveconf:batchid}
107 AND (liked=1 OR song_end_type=0)
108 GROUP BY song_id
109 ORDER BY duration DESC
110 LIMIT 10;
```

```
113 CREATE TABLE IF NOT EXISTS top_10_unsubscribed_users
114 (
115 user_id STRING,
116 duration INT
117 )
118 PARTITIONED BY (batchid INT)
119 ROW FORMAT DELIMITED
120 FIELDS TERMINATED BY ','
121 STORED AS TEXTFILE;
122
123 INSERT OVERWRITE TABLE top_10_unsubscribed_users
124 PARTITION(batchid=${hiveconf:batchid})
125 SELECT
126 ed.user_id,
127 SUM(ABS(CAST(ed.end_ts AS DECIMAL(20,0))-CAST(ed.start_ts AS DECIMAL(20,0)))) AS duration
128 FROM enriched_data ed
129 LEFT OUTER JOIN subscribed_users su
130 ON ed.user_id=su.user_id
131 WHERE ed.status='pass'
132 AND ed.batchid=${hiveconf:batchid}
133 AND (su.user_id IS NULL OR (CAST(ed.timestamp AS DECIMAL(20,0)) > CAST(su.subscn_end_dt AS DECIMAL(20,0))))
134 GROUP BY ed.user_id
135 ORDER BY duration DESC
136 LIMIT 10;
```

Run script: `./data_analysis.sh`

```
[acadgild@localhost music]$ sh data_analysis.sh
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12.jar]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]

Logging initialized using configuration in jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/hive-common-2.3.2.jar!/hive-log4j.properties
OK
Time taken: 17.679 seconds
OK
Time taken: 1.786 seconds
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff
Hive 1.X releases.
Query ID = acadgild_20190123051101_2c4da4e6-9bb6-4a5f-902b-aa88c0afac2d
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1548195624342_0005, Tracking URL = http://localhost:8088/proxy/application_1548195624342_0005/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1548195624342_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2019-01-23 05:11:40,526 Stage-1 map = 0%, reduce = 0%
2019-01-23 05:12:09,540 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 10.49 sec
2019-01-23 05:12:36,903 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 17.52 sec
MapReduce Total cumulative CPU time: 17 seconds 520 msec
Ended Job = job_1548195624342_0005
Launching Job 2 out of 2
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
```


BIG DATA PROJECT

The tables have been created in the Hive,

```
hive> show tables;
OK
connected_artists
enriched_data
formatted_input
song_artist_map
station_geo_map
subscribed_users
top_10_royalty_songs
top_10_stations
top_10_unsubscribed_users
users_artists
users_behaviour
Time taken: 1.293 seconds, Fetched: 11 row(s)
hive> █
```

Problem statement 1:

Determine top 10 **station_id(s)** where maximum number of songs were played, which were liked by unique users.

*select * from top_10_stations;*

```
hive> select * from top_10_stations;
OK
ST413    7      8      1
ST402    6      7      1
ST408    6      7      1
ST410    5      5      1
ST403    5      6      1
ST412    4      5      1
ST411    3      3      1
ST401    3      4      1
ST406    3      3      1
ST409    3      6      1
Time taken: 7.164 seconds, Fetched: 10 row(s)
hive> █
```

Problem statement 2:

Determine total duration of songs played by each type of user, where type of user can be 'subscribed' or 'unsubscribed'. An unsubscribed user is the one whose record is either not present in Subscribed_users lookup table or has subscription_end_date earlier than the timestamp of the song played by him.

*select * from users_behaviour;*

```
hive> select * from users_behaviour;
OK
SUBSCRIBED      1215318357      1
UNSUBSCRIBED    1061291483      1
Time taken: 0.694 seconds, Fetched: 2 row(s)
hive> █
```

Problem statement 3:

Determine top 10 connected artists. Connected artists are those whose songs are most listened by the unique users who follow them.

*select * from connected_artists;*

```
hive> select * from connected_artists;
OK
A302      10      1
A301      10      1
A303       4      1
A304       3      1
A305       2      1
A300       2      1
Time taken: 0.622 seconds, Fetched: 6 row(s)
hive> █
```

Problem statement 4:

Determine top 10 songs which generated the maximum revenue. Royalty applies to a song only if it was liked or was completed successfully or both.

*select * from top_10_royalty_songs;*

```
hive> select * from top_10_royalty_songs;
OK
S201      182809906      1
S204      143716685      1
S202      124279939      1
S206      108971273      1
S203      101311339      1
S207       97636973      1
S208       96202673      1
S209       88707006      1
S200       67514012      1
S205       27758921      1
Time taken: 0.723 seconds, Fetched: 10 row(s)
hive> █
```

Problem statement 5:

Determine top **10 unsubscribed** users who listened to the songs for the longest duration

*select * from top_10_unsubscribed_users;*

```
hive> select * from top_10_unsubscribed_users;
OK
U110      37745000
U114      35601000
U107      27450068
U109      26789451
U115      19754610
U111      18456248
U103      15470012
U112      12100349
U106      10078523
U107      9870113
Time taken: 0.493 seconds, Fetched: 10 row(s)
hive> █
```

Now, we need to export all the data to the **MYSQL** using sqoop, run the script **data_export.sh**

4.5 Stage – 5 – Data Storage in MYSQL

Using the bash file shown below, **data_export.sh** we are going to export the data from the HIVE tables into MYSQL using **SQOOP** export.

```
1  #!/bin/bash
2
3  #This script is not working.
4  #Either change table to text or use STRING as type of partitioned column
5
6  batchid=`cat /home/acadgild/project/logs/current-batch.txt`
7  LOGFILE=/home/acadgild/project/logs/log_batch_$batchid
8
9  echo "Creating mysql tables if not present..." >> $LOGFILE
10
11 mysql < /home/acadgild/project/scripts/create_schema.sql
12
13 echo "Running sqoop job for data export..." >> $LOGFILE
14
15 sqoop export --connect jdbc:mysql://localhost/project --username root --password acadgild --table top_10_stations --export-dir
16 hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_stations/batchid=$batchid --input-fields-terminated-by ',' -m 1
17
18 sqoop export --connect jdbc:mysql://localhost/project --username root --password acadgild --table users_behaviour --export-dir
19 hdfs://localhost:9000/user/hive/warehouse/project.db/users_behaviour/batchid=$batchid --input-fields-terminated-by ',' -m 1
20
21 sqoop export --connect jdbc:mysql://localhost/project --username root --password acadgild --table connected_artists --export-dir
22 hdfs://localhost:9000/user/hive/warehouse/project.db/connected_artists/batchid=$batchid --input-fields-terminated-by ',' -m 1
23
24 sqoop export --connect jdbc:mysql://localhost/project --username root --password acadgild --table top_10_royalty_songs --export-dir
25 hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_royalty_songs/batchid=$batchid --input-fields-terminated-by ',' -m 1
26
27 sqoop export --connect jdbc:mysql://localhost/project --username root --password acadgild --table top_10_unsubscribed_users --export-dir
28 hdfs://localhost:9000/user/hive/warehouse/project.db/top_10_unsubscribed_users/batchid=$batchid --input-fields-terminated-by ',' -m 1
```

BIG DATA PROJECT

create_schema.sql – The below schema will create the database and tables in the MySQL.

```
1 CREATE DATABASE IF NOT EXISTS project;
2
3 USE project;
4
5 CREATE TABLE IF NOT EXISTS top_10_stations
6 (
7     station_id VARCHAR(50),
8     total_distinct_songs_played INT,
9     distinct_user_count INT
10 );
11
12 CREATE TABLE IF NOT EXISTS users_behaviour
13 (
14     user_type VARCHAR(50),
15     duration BIGINT
16 );
17
18 CREATE TABLE IF NOT EXISTS connected_artists
19 (
20     artist_id VARCHAR(50),
21     user_count INT
22 );
23
24 CREATE TABLE IF NOT EXISTS top_10_royalty_songs
25 (
26     song_id VARCHAR(50),
27     duration BIGINT
28 );
29
30 CREATE TABLE IF NOT EXISTS top_10_unsubscribed_users
31 (
32     user_id VARCHAR(50),
33     duration BIGINT
34 );
35
36 commit;
```

The database *project* has been exported from **HIVE** and the below screenshot shows the exported data to **MYSQL**.

```
mysql> show tables;
+-----+
| Tables_in_project |
+-----+
| connected_artists |
| top_10_royalty_songs |
| top_10_stations |
| top_10_unsubscribed_users |
| users_behaviour |
+-----+
5 rows in set (0.01 sec)

mysql> █
```

```
mysql> select * from top_10_stations;
```

station_id	total_distinct_songs_played	distinct_user_count
ST413	7	8
ST402	6	7
ST408	6	7
ST410	5	5
ST403	5	6
ST412	4	5
ST411	3	3
ST401	3	4
ST406	3	3
ST409	3	6

```
10 rows in set (0.00 sec)

mysql>
```

```
mysql> select * from users_behaviour;
```

user_type	duration
SUBSCRIBED	1215318357
UNSUBSCRIBED	1061291483

```
2 rows in set (0.01 sec)

mysql> select * from connected_artists;
```

artist_id	user_count
A302	10
A301	10
A303	4
A304	3
A305	2
A300	2

```
6 rows in set (0.00 sec)

mysql>
```

```
mysql> select * from top_10_royalty_songs;
```

song_id	duration
S201	182809906
S204	143716685
S202	124279939
S206	108971273
S203	101311339
S207	97636973
S208	96202673
S209	88707006
S200	67514012
S205	27758921

```
10 rows in set (0.01 sec)

mysql> select * from top_10_unsubscribed_users;
```

user_id	duration
U110	37745000
U114	35601000
U107	27450068
U109	26789451
U115	19754610
U111	18456248
U103	15470012
U112	12100349
U106	10078523
U107	9870113

```
10 rows in set (0.00 sec)

mysql>
```

BIG DATA PROJECT

JOB SCHEDULING -

Now after exporting data into MySQL **batchid** will be incremented to additional 1, so one batch of data operation is successfully completed and new batch of data will be loaded for the analysis after every 3 hours.

```
sh /home/acadgild/project/scripts/data_export.sh  
echo "Incrementing batchid..." >> $LOGFILE  
batchid=`expr $batchid + 1`  
echo -n $batchid > /home/acadgild/project/logs/current-batch.txt
```

We can check logs to track the behavior of the operations we have done on the data and overcome failures in the pipeline and we can see the **batchid** incremented value in **current-batch.txt**

```
[acadgild@localhost logs]$ ls -l  
total 12  
-rwxrwxr-x. 1 acadgild acadgild  1 Jan 29 07:00 current-batch.txt  
-rw-rw-r--. 1 acadgild acadgild 1222 Jan 29 06:04 log_batch_1  
-rw-rw-r--. 1 acadgild acadgild  399 Jan 29 06:30 log_batch_1???  
[acadgild@localhost logs]$ cat current-batch.txt  
2[acadgild@localhost logs]$
```

The log file captured all the data and steps we performed so far,

```
[acadgild@localhost logs]$ cat log_batch_1  
Starting daemons  
Creating LookUp Tables  
Populating LookUp Tables  
Creating hive tables on top of hbase tables for data enrichment and filtering...  
Placing data files from local to HDFS...  
Running pig script for data formatting...  
Running hive script for formatted data load...  
Running hive script for data enrichment and filtering...  
Copying valid and invalid records in local file system...  
Deleting older valid and invalid records from local file system...  
Running hive script for data analysis...  
Incrementing batchid...  
[acadgild@localhost logs]$
```

BIG DATA PROJECT

Wrapping all the scripts inside the single script file and scheduling this file to run at the periodic interval of every 3 hours.

```
1 #!/bin/bash
2
3 python /home/acadgild/project/scripts/generate_web_data.py
4
5 python /home/acadgild/project/scripts/generate_mob_data.py
6
7 sh /home/acadgild/project/scripts/start-daemons.sh
8
9 sh /home/acadgild/project/scripts/populate-lookup.sh
10
11 sh /home/acadgild/project/scripts/dataformatting.sh
12
13 sh /home/acadgild/project/scripts/data_enrichment.sh
14
15 sh /home/acadgild/project/scripts/data_analysis.sh
```

The wrapper.sh will be running for every 3 hours as per the job scheduling done below, as per the above order the wrapper.sh will run the scripts.

Creating **Crontab** to schedule the wrapper.sh script to run for every 3 hour interval -

```
#Run every 3 hours
* */3 * * * date>>/home/acadgild/project/scripts/wrapper.sh >> /home/acadgild/project/scripts/jobscheduling.log

~
~
~
~
~
~
~
~
~
~
```

The **crontab** job scheduler will run the **wrapp.sh** every 3 hours and for every 3 hours we will get incremental batch ID's. As per the requirement the job scheduling is done.

Highlights of the Project :

- LookUp tables in HBASE have been integrated with actual flow of data.
- Joins were optimized for analysis. Data was enriched with new fields and using broadcast maps on Lookup tables so as to avoid joins.
- Data cleaning, validation, enrichment and analysis have been automated using bash scripts and schedulers.
- Logs have been maintained to track the behavior and overcome failures in the pipeline.

Project End Conclusion:

All the data operation has been performed as per the sequence mentioned in the **wrapper.sh** file and obtained results successfully for the leading music company to make appropriate business strategies. The results can be used by data science or machine learning pipelines for further forecast and form visualization on the analyzed data.

--- END ---